1. What do you think applying this filter to a grayscale image will do?

$$\begin{bmatrix} 0 & 1 & -1 & 0 \\ 1 & 3 & -3 & -1 \\ 1 & 3 & -3 & -1 \\ 0 & 1 & -1 & 0 \end{bmatrix}$$

- ○ Detect image contrast
- ● Detect vertical edges
- ○ Detect horizontal edges
- ○ Detect 45 degree edges

2. Suppose your input is a 300 by 300 color (RGB) image, and you are not using a convolutional network. If the first hidden layer has 100 neurons, each one fully connected to the input, how many parameters does this hidden layer have (including the bias parameters)?

- ○ 9,000,001
- ○ 9,000,100
- ○ 27,000,001
- ● 27,000,100

3. Suppose your input is a 300 by 300 color (RGB) image, and you use a convolutional layer with 100 filters that are each 5x5. How many parameters does this hidden layer have (including the bias parameters)?

- ○ 2501
- ○ 2600
- ○ 7500
- ● 7600

4. You have an input volume that is 63x63x16, and convolve it with 32 filters that are each 7x7, using a stride of 2 and no padding. What is the output volume?

- ● 29x29x32
- ○ 16x16x32
- ○ 16x16x16
- ○ 29x29x16

5. You have an input volume that is 15x15x8, and pad it using "pad=2." What is the dimension of the resulting volume (after padding)? `1 point`

- ● 19x19x8
- ○ 19x19x12
- ○ 17x17x10
- ○ 17x17x8

6. You have an input volume that is 63x63x16, and convolve it with 32 filters that are each 7x7, and stride of 1. You want to use a "same" convolution. What is the padding? `1 point`

- ○ 1
- ○ 2
- ● 3
- ○ 7

7. You have an input volume that is 32x32x16, and apply max pooling with a stride of 2 and a filter size of 2. What is the output volume? `1 point`

- ○ 15x15x16
- ● 16x16x16
- ○ 16x16x8
- ○ 32x32x8

8. Because pooling layers do not have parameters, they do not affect the backpropagation (derivatives) calculation. `1 point`

- ○ True
- ● False

9. In lecture we talked about "parameter sharing" as a benefit of using convolutional networks. Which of the following statements about parameter sharing in ConvNets are true? (Check all that apply.)

- ☐ It allows gradient descent to set many of the parameters to zero, thus making the connections sparse.

- ☑ It allows a feature detector to be used in multiple locations throughout the whole input image/input volume.

- ☐ It allows parameters learned for one task to be shared even for a different task (transfer learning).

- ☑ It reduces the total number of parameters, thus reducing overfitting.

10. In lecture we talked about "sparsity of connections" as a benefit of using convolutional layers. What does this mean?

- ⭘ Regularization causes gradient descent to set many of the parameters to zero.

- ⦿ Each activation in the next layer depends on only a small number of activations from the previous layer.

- ⭘ Each layer in a convolutional network is connected only to two other layers

- ⭘ Each filter is connected to every channel in the previous layer.